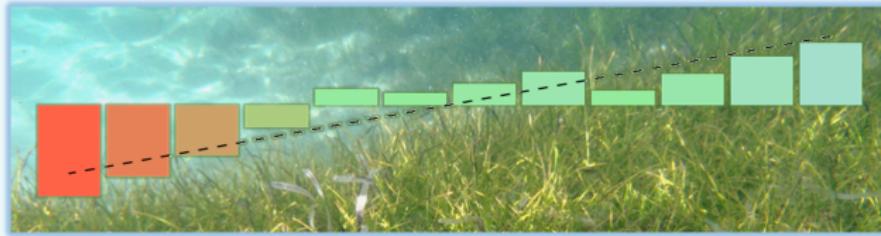


SWMPPr: An R package for estuarine water quality time series

Marcus W. Beck

ORISE, USEPA NHEERL Gulf Ecology Division
Email: beck.marcus@epa.gov

April 15, 2015



Overview

A mixture of things...

- What is NERRS/SWMP and motivation for creating the package
- The process of package development (briefly)
- What can SWMPr do
- What has SWMPr done

What is NERRS/SWMP?

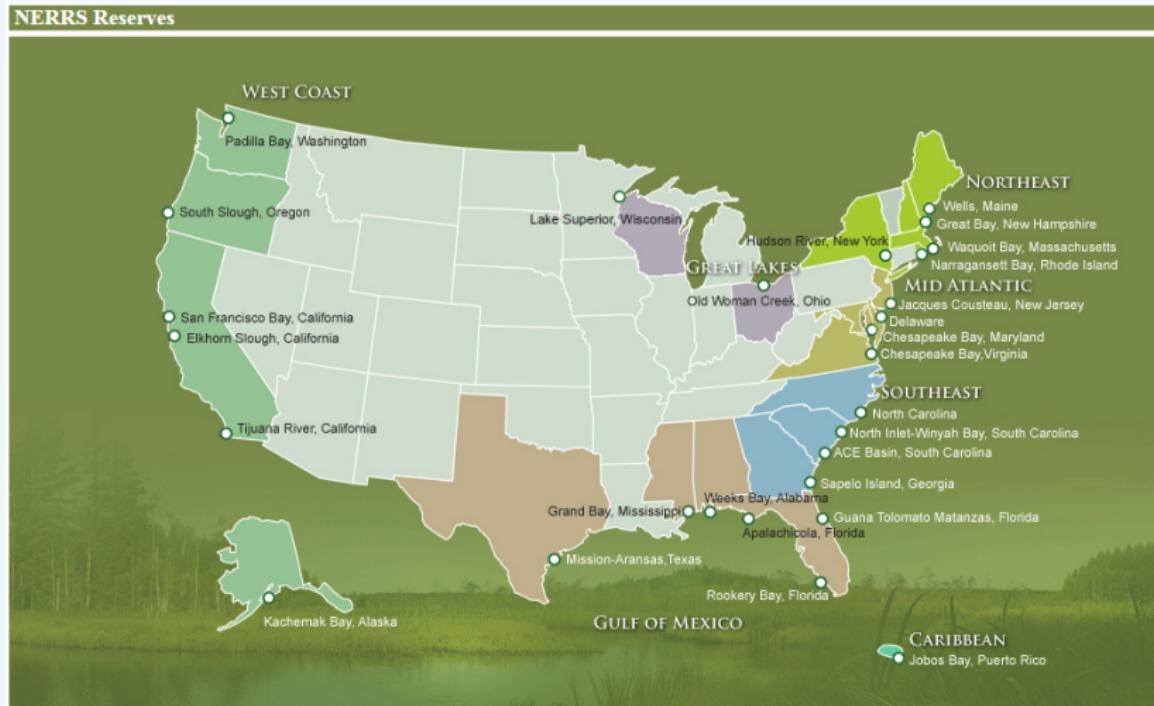
NERRS

National Estuarine Research Reserve System, established by Coastal Zone Management Act of 1972. Focus on *long-term research, monitoring, education, and stewardship* for more effective coastal management.

SWMP

System Wide Monitoring Program, initiated in 1995 to provide *continuous monitoring* data at over 300 stations in each of the 28 NERRS reserves

What is NERRS/SWMP?



<http://nerrs.noaa.gov/ReservesMap.aspx>

What is NERRS/SWMP?

Each reserve has fixed, continuous monitoring stations for ***water quality*** (15 min), ***meteorology*** (15 min), and ***nutrients*** (monthly)

The parameters for a station are specific to the parameter type

Water quality

temp, spcond, sal,
do_pct, do_mgl,
depth, cdepth, level,
clevel, ph, turb,
chlfluor

Meteorology

atemp, rh, bp, wspd,
maxwspd, wdir,
sdwdir, totpar,
totprcp, cumprcp,
totsorad

Nutrients

po4f, chla_n, no3f,
no2f, nh4f, no23f,
ke_n, urea

What is NERRS/SWMP?

Data maintained by the Centralized Data Management Office (CDMO)

The screenshot shows the homepage of the NERRS/SWMP website. At the top, there is a navigation bar with links for Home, About CDMO, About Data, Get Data, Web Services, and Contact CDMO. Below the navigation bar is a large banner image featuring a white heron standing in a wetland area with a river and fields in the background.

The main content area is divided into three columns:

- View / Download Data**: This column features a large map of the United States with various monitoring sites marked. Overlaid on the map is the text "View / Download Data". Below the map is a link to "Requested Citation Format".
- Real Time Monitoring Data**: This column includes a dropdown menu labeled "Choose Reserve..." with options "GTMPCMET 10/08/14 09:45 AM" and "GTMPCWQ 10/08/14 09:45 AM". It also shows a photograph of a weather station at a research site. Below the image are real-time data readings:
 - Air Temperature: 27.8 °C (82 °F)
 - Wind Speed: 1.1 m/Sec (02 mph)
 - Water Temperature: 22.7 °C (73 °F)
 - Salinity: 7.1 PPT
 - Dissolved Oxygen: 4.7 mg/L
- CDMO News**: This column contains a news announcement about the launch of the SWMP Mobile application. It includes a link to the mobile application's website: www.nerrsdata.org/mobile. Below this, there is another news item about the updated Data Export System, which now includes enhanced graphing capabilities. It also includes a link to the Data Export System: [Data Export System!](#).

What is NERRS/SWMP?

As of April 10, > 58 million SWMP data records available from CDMO

Raw data will look like this...

A	B	C	D	E	F	G	H	I	J	K	L	
1	StationCode	isSWMP	DateTimeStamp	Historical	Provisional	CollMethod	REP	F_Record	PO4F	F_PO4F	NH4F	F_NH4F
2	apacpnut	P	1/10/2012 10:20	0	1	1	1		0.003 <-4> [SBL]		0.03	<0>
3	apacpnut	P	2/7/2012 11:41	0	1	1	1		0.005 <0>		0.019	<0>
4	apacpnut	P	3/5/2012 11:51	0	1	1	1		0.003 <-4> [SBL]		0.041	<0>
5	apacpnut	P	4/4/2012 10:30	0	1	1	1		0.003 <-4> [SBL]		0.043	<0>
6	apacpnut	P	5/9/2012 10:12	0	1	1	1		0.003 <0>		0.053	<0>
7	apacpnut	P	5/9/2012 10:15	0	1	1	2		0.003 <-4> [SBL]		0.022	<0>
8	apacpnut	P	5/9/2012 10:20	0	1	1	3		0.003 <0>		0.016	<0>
9	apacpnut	P	6/5/2012 8:30	0	1	1	1		0.003 <-4> [SBL]		0.04	<0>
10	apacpnut	P	7/3/2012 9:58	0	1	1	1 {CSM}		0.004 <0>		0.094	<0>
11	apacpnut	P	7/3/2012 9:59	0	1	1	2 {CSM}		0.004 <0>		0.066	<0>
12	apacpnut	P	7/3/2012 10:01	0	1	1	3 {CSM}		0.005 <0>		0.069	<0>
13	apacpnut	P	8/7/2012 9:53	0	1	1	1 {CSM}		0.003 <-4> [SBL]		0.05	<0>
14	apacpnut	P	9/5/2012 10:56	0	1	1	1		0.003 <-4> [SBL]		0.026	<0>
15	apacpnut	P	10/2/2012 9:22	0	1	1	1		0.003 <-4> [SBL]		0.042	<0>
16	apacpnut	P	10/2/2012 9:27	0	1	1	2		0.003 <-4> [SBL]		0.024	<0>
17	apacpnut	P	10/2/2012 9:32	0	1	1	3		0.003 <0>		0.042	<0>
18	apacpnut	P	11/6/2012 10:30	0	1	1	1		0.003 <-4> [SBL]		0.07	<0>
19	apacpnut	P	11/26/2012 11:39	0	1	1	1		0.003 <-4> [SBL]		0.041	<0>

What is the problem?

An invaluable data source but no recent comparative analyses between systems

NERRS researchers, managers, and technicians need more tools for trend analysis

Some specific issues:

- Knowing what data to use and how to obtain
- Dealing with QAQC columns or removing ‘bad’ observations
- Combining data for comparison
- Issues inherent with time series, e.g., signal vs. noise, data quantity

What is the (potential) solution?



SWMPPr v2.0.0 is now available on CRAN!

```
> install.packages('SWMPPr')  
> library(SWMPPr)
```

Currently no vignette, but working on a manuscript

The process of package development

The development version lives on GitHub:

<https://github.com/fawda123/SWMPr>

The package development process was much simplified using RStudio and the Hadleyverse (specifically `devtools`, `roxygen2`)

In RStudio, create a package template:

`File > New Project > New Directory > R Package`, with options for Git version control

Package does not have to be on CRAN to distribute...

The process of package development

Follow the advice here: <http://r-pkgs.had.co.nz/>

R packages by Hadley Wickham

[Table of contents ▾](#)

Want to learn from me in person?
I'm next teaching in Chicago, May
27-28.

Want a physical copy of this
material? [Buy from amazon!](#).

Contents

How to contribute

[Edit this page](#)

R packages

This is the in-progress book site for "**R packages**". It will be published with O'Reilly around March 2015. You can [pre-order](#) a copy from amazon.



Packages are the fundamental units of reproducible R code. They include reusable R functions, the documentation that describes how to use them, and sample data. In this section you'll learn how to turn your code into packages that others can easily download and use. Writing a package can seem overwhelming at first. So start with the basics and improve it over time. It doesn't matter if your first version isn't perfect as long as the next version is better.

Getting started

- [Introduction](#)
- [Package structure](#)

What can SWMPr do?

SWMPr functions are grouped into three categories that describe their use in the ‘data workflow’

Retrieve

all_params
all_params_dtrng
import_local
import_remote
single_param
site_codes
site_codes_ind

Organize

comb
qaqc
qaqcchk
rem_reps
setstep
subset

Analyze

aggreswmp
aggremetab
ecometab
decomp
decomp_cj
hist
lines
na.approx
plot
plot_metab
plot_summary
smoother

How are data *retrieved*?

The first challenge is to determine the station, parameter, and date range of interest - Check the available data on the CDMO website

Also familiarize yourself with the NERRS/SWMP naming convention

Site (reserve), ***station***, and ***parameter type*** are identified by a 7 or 8 character name

E.g., elkcwmet

- elk: site, Elkhorn Slough
- cw: station, Caspian Weather Station
- met: parameter type (weather)

How are data *retrieved*?

SWMP data can be imported from CDMO into R three ways

1) From a local path: `import_local`

Advantages:

- Ideal for large datasets
- Most recent data

Disadvantages:

- Outside of R
- Only works for one type of data request from CDMO

How are data *retrieved*?

SWMP data can be imported from CDMO into R three ways

2) retrieve SWMP data from a third-party server: `import_remote`

Advantages:

- Fast import!
- No requests to CDMO

Disadvantages:

- Data are not current
- Data have already been minimally processed

How are data *retrieved*?

SWMP data can be imported from CDMO into R three ways

- 3) Use the existing CDMO web services to import directly:
`single_param`, `all_params`, `all_params_dtrng`

Advantages:

- Data will always be current
- Customized requests

Disadvantages:

- IP address must be registered
- Quantity severely limited

How are data *retrieved*?

The end result is the same - data are imported as a `swmpr` S3 object

```
> dat <- import_remote('kacsswq')
> class(dat)

## [1] "swmpr"      "data.frame"

> head(dat, 1)

##   datetimestamp temp spcond sal do_pct do_mgl depth cdepth
## 1 2004-01-01     2     42  26    101     12    0.7     NA
##   level clevel ph turb chlfluor
## 1    NA      NA  8     6       NA

> names(attributes(dat))

## [1] "names"      "row.names"   "class"      "station"
## [5] "parameters" "qaqc_cols"   "date_rng"   "timezone"
## [9] "stamp_class"
```

How are data *retrieved*?

The remaining functions have `swmpr` methods

```
> methods(class = 'swmpr')

## [1] aggregemetab.swmpr*    aggreswmp.swmpr*
## [3] comb.swmpr*           decomp.swmpr*
## [5] decomp_cj.swmpr*      ecometab.swmpr*
## [7] hist.swmpr*           lines.swmpr*
## [9] na.approx.swmpr*      plot.swmpr*
## [11] plot_metab.swmpr*     plot_summary.swmpr*
## [13] qaqc.swmpr*           qaqcchk.swmpr*
## [15] rem_reps.swmpr*       setstep.swmpr*
## [17] smoother.swmpr*       subset.swmpr*
##
## Non-visible functions are asterisked
```

`swmpr` objects also inherit methods from the `data.frame` class

How are data *organized*?

Data organization depends on the analysis needs - it is neither fun nor straightforward (common opinion, not mine)

What are some challenges?

- Imported data have QAQC columns
- Extra columns/rows
- Maybe we don't care about all the parameters
- Data from separate sites are in separate objects

The *organize* functions are specific to the SWMP data but many of the principles apply to generic time series

How are data *organized*?

A relevant example - we want to compare time series from different sites

- Data may have arbitrary time steps that do not match between sites
- Date ranges may also differ

The `setstep` and `comb` functions address these issues!

```
> met <- import_remote('apaebmet')
> wq <- import_remote('apacpwq')
> dat <- comb(met, wq) # tada!
```

How are data *organized*?

The `setstep` function is used within `comb` to standardize the time steps for each input object

A tricky problem - actual observations which may occur on an arbitrary step must be matched to a set time step

This function uses ‘fast-ordered joins’ from the `data.table` package using the ‘nearest’ method

Also must define a threshold for matching: +/- some buffer of allowance beyond which matches are discarded

How are data *organized*?

Mechanistically, `setstep` does the following for each data object:

- Creates a continuous ‘master’ time series at defined step using first/last time stamps
- Matches existing observations to standardized using ‘nearest’ join method
- Calculates difference in time between matched and standardized step, discards those beyond threshold

Data objects, now standarized, are then combined by absolute matching of time steps, again using `data.table`

How are data *organized*?

```
> dim(met)
## [1] 490847      11

> dim(wq)
## [1] 455808      13

> # standardize time step to two hours
> # maximum difference for matching 30 minutes
> # combine only overlapping time ranges
> dat <- comb(wq, met, timestep = 120, differ = 30,
+   method = 'intersect')
> dim(dat)
## [1] 56977      23
```

How are data *organized*?

```
> head(dat, 4)
```

```
##           datetimestamp atemp  rh    bp wspd maxwspd wdir
## 1 2001-12-31 23:00:00      4 69 1017      4       NA 347
## 2 2002-01-01 01:00:00      3 75 1017      3       NA   9
## 3 2002-01-01 03:00:00      2 77 1018      3       NA 331
## 4 2002-01-01 05:00:00      1 82 1019      4       NA   0
##   sdwdir totpar totprcp totsorad temp spcond sal do_pct
## 1     NA     0      NA      NA     NA     NA     NA     NA
## 2     NA     0      NA      NA     12     37     24    104
## 3     NA     0      NA      NA     12     40     26     99
## 4     NA     0      NA      NA     11     42     26     98
##   do_mgl depth cdepth level clevel ph turb chlfluor
## 1     NA     NA     NA     NA     NA    NA     NA     NA
## 2     10     2     NA     NA     NA    NA     3     NA
## 3      9     2     NA     NA     NA    NA     4     NA
## 4      9     2     NA     NA     NA    NA     5     NA
```

How are data *analyzed*?

Time series analysis can range from very general to very specific

SWMPr functions include...

General

- `na.approx`
- `smoother`
- `aggreswmp`
- `plot`
- `hist`
- `lines`

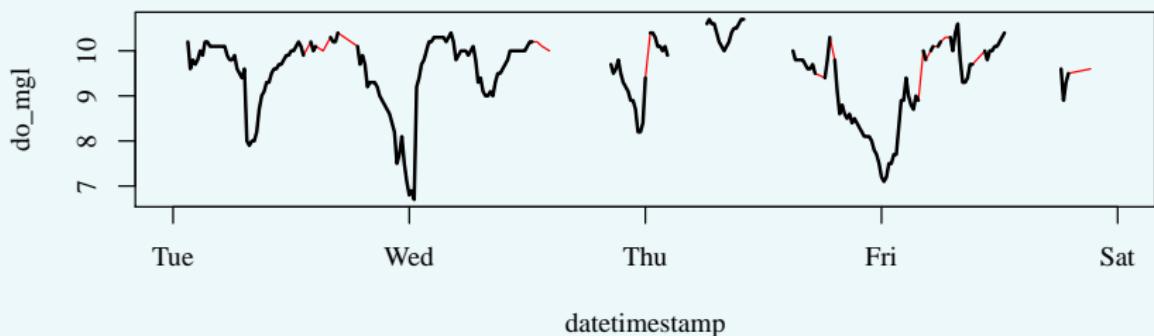
Specific

- `decomp`
- `decomp_cj`
- `eometab`
- `aggremetab`
- `plot_metab`
- `plot_summary`

How are data *analyzed*?

Some examples... fill missing data with `na.approx`

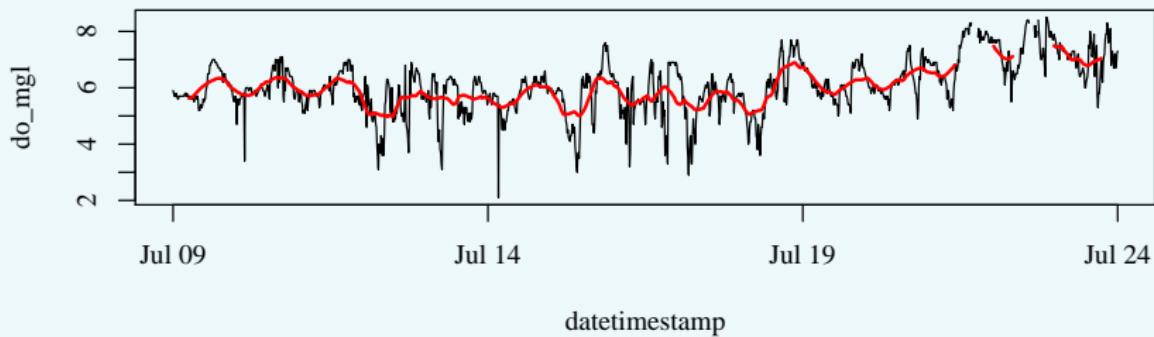
```
> # import, organize
> data(apadbwq)
> dat <- qaqc(apadbwq)
> dat <- subset(dat, select = 'do_mgl',
+   subset = c('2013-01-22 00:00', '2013-01-26 00:00')
+ )
>
> # interpolate, plot
> filled <- na.approx(dat, params = 'do_mgl', maxgap = 10)
> plot(filled, col = 'red'); lines(dat, lwd = 2)
```



How are data *analyzed*?

Some examples... smooth data with `smoother`

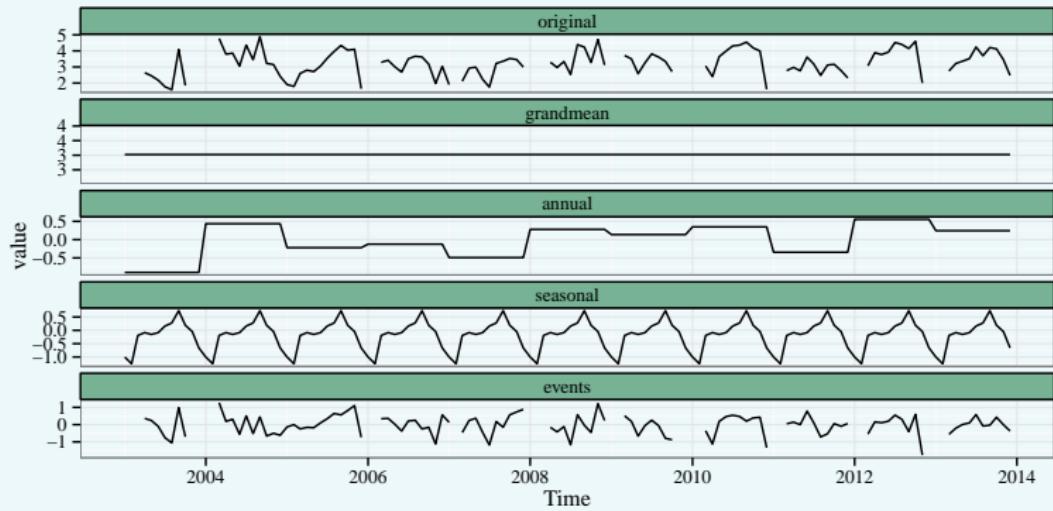
```
> # import, organize
> data(apadbwq)
> dat <- qaqc(apadbwq)
> dat <- subset(dat, select = 'do_mgl',
+   subset = c('2012-07-09 00:00', '2012-07-24 00:00')
+ )
>
> # smooth, plot
> dat_smooth <- smoother(dat, window = 50, params = 'do_mgl')
> plot(dat); lines(dat_smooth, col = 'red', lwd = 2)
```



How are data *analyzed*?

Some examples... additive decomposition with `decomp_cj`

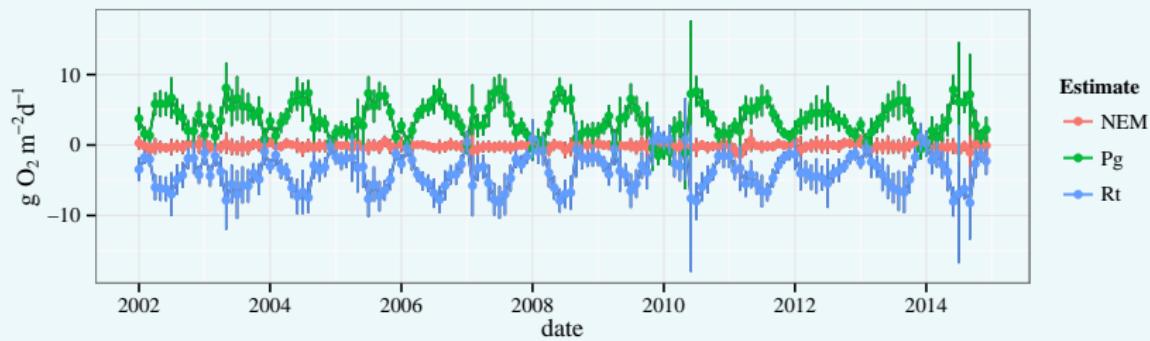
```
> # import, organize
> dat <- import_remote('cbmocnut')
> dat$chla_n <- log(dat$chla_n)
>
> # additive decompsision of chl, annual
> decomp_cj(dat, 'chla_n', type = 'add')
```



How are data *analyzed*?

Some examples... estimate metabolism with **ecometab**

```
> ## import water quality and weather data, combine
> wq <- import_remote('apebbwq')
> met <- import_remote('apaebmet')
> dat <- comb(wq, met)
>
> ## metabolism, plot
> res <- ecometab(dat, metab_units = 'grams')
> plot_metab(res, by = 'months')
```



SWMPPr applications

The most common question - has there been a change over time at my site?

The functions in the package can help answer this very general question...

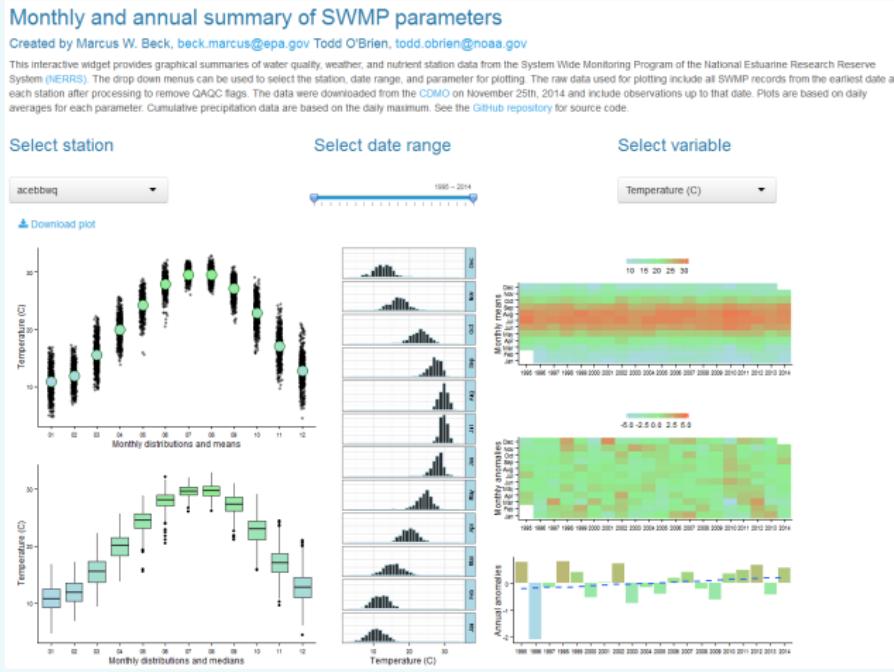
...but it's often easier to interactively evaluate the data!

Two shiny applications, hosted on shinyapps.io, allow users to visualize trends in SWMP data

These apps were created with SWMPPr functions or use them ‘reactively’

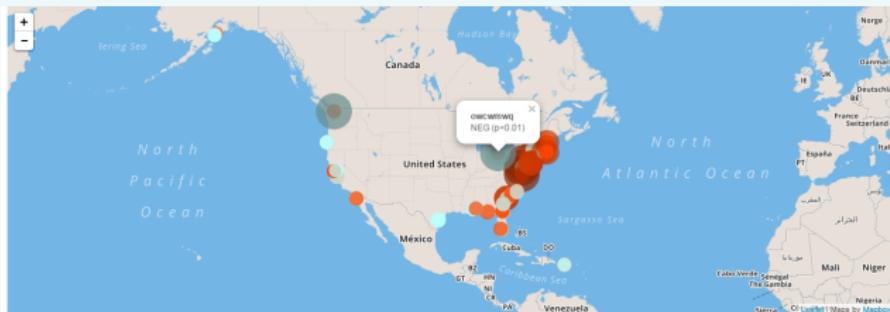
SWMPr applications

SWMP summary plots: https://beckmw.shinyapps.io/swmp_summary/



SWMPr applications

SWMP trends map: https://beckmw.shinyapps.io/swmp_comp/



Trends in SWMP parameters

Created by Marcus W. Beck, beck.marcus@epa.gov, Todd O'Brien, todd.obrien@noaa.gov

This widget is an interactive tool to evaluate trends in SWMP data. Trends are described by an increase or decrease in values over time using a simple linear regression of summarized data. The regression for each station can be viewed by clicking on a map location. Trends at each station are plotted as circles that identify the direction and significance of the trend. The trend direction is blue for decreasing and red for increasing. The significance is indicated by radius of the circle and color shading where larger points with darker colors indicate a strong trend. Original data are available from <http://swmp.noaa.gov/>. The map is centered at 34.44, -93.96 with a zoom level of 3.

Select parameter:

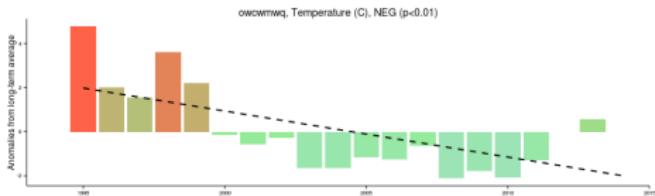
wq: Temperature (C)

Summarize by:

Years: anomalies

Select date range:

1964 - 2014



Conclusions

The SWMPr package provides an R-centric approach to *retrieve*, *organize*, and *analyze* estuary data

SWMPr is meant to *augment*, not replace, existing data management programs (i.e., CDMO web services)

Benefits of the package: deals with lots of the heavy lifting with large, unrefined datasets

Some points of concern: do not use functions ad nauseum, always refer to help files and optional arguments

Just because you can do something doesn't mean you should!

Conclusions



Visit the development repo for the most recent version or to submit a bug report: <https://github.com/fawda123/SWMPPr>

To get this presentation: https://github.com/fawda123/swmpr_pres

To learn more about NERRS: <http://www.nerrs.noaa.gov/>

To learn more about SWMP/CDMO: <http://cdmo.baruch.sc.edu/>